

AI-Based Intelligent Retrieval and Quality-Aware Curation of Scientific Microscopy Images

Pranet Pallati¹, Gollakota Charan Deep¹, Pooja Vunnam¹, and Ms. C. Bhavana²

¹Department of Computer Science and Engineering, Vardhaman College of Engineering, Hyderabad, Telangana, India
Emails: {24881A05B7, 24881A0586, 24881A05B5}@student.vardhaman.org

²Assistant Professor, Department of Computer Science and Engineering, Vardhaman College of Engineering, bhavana1817@vardhaman.org

Abstract—Scientific microscopy repositories are expanding rapidly across material science and biology, yet retrieval and curation remain severely constrained by acquisition heterogeneity (varying accelerating voltages, detectors, and magnifications) and unquantified quality risks. In this work, we present SCI-INTEL, a reproducible platform integrating metadata-aware ingestion, acquisition-aware representation adaptation, controlled artifact screening, spatial localization, uncertainty-aware abstention, and deterministic comparative evidence aggregation. Evaluating on 6,085 scientific micrographs (774 HCCI SEM, 4,591 Carinthia SEM, and 720 BBBC021 optical micrographs), we demonstrate that lightweight representation adaptation reduces the observed acquisition-geometry similarity gap by 66.23% (from 0.2016 to 0.0681, $p = 5.03e-36$, Cohen's $d_z = 2.19$) while achieving Recall@5 of 0.9921 and MRR of 0.5261 under realistic unmasked distractor retrieval (Protocol U). Our evaluation reveals an empirical specialization trade-off: while adapted representations optimize cross-instrument retrieval, frozen DINOv2 ViT-S/14 features retain superior sensitivity to controlled synthetic artifacts (Macro F1 = 0.6837 vs. 0.6323). SCI-INTEL resolves this trade-off via a modular dual-representation architecture that routes each representation to its specialized task without learned fusion. Uncertainty-aware abstention safely routes low-confidence cases (confidence < 0.40 or normalized entropy > 0.75) to specialist review, while the evidence layer achieves 100.0% valid comparative retrieval across an evaluated $N=55$ query cohort. Our evaluation does not claim unseen-specimen generalization, physical defect confirmation, clinical diagnosis, or human interpretation improvement.

Index Terms—Scientific Image Retrieval, Microscopy, Representation Learning, Acquisition Robustness, Image Quality Assessment, Anomaly Screening, Uncertainty Estimation, Scientific Data Curation.

I. INTRODUCTION

Modern scientific microscopy imaging has evolved into an essential foundation for characterization in materials science, crystallography, and biology [12, 14]. Instruments such as scanning electron microscopes (SEM), transmission electron microscopes (TEM), and focused ion beam systems routinely generate millions of high-resolution micrographs. However, retrieving and organizing these images across large-scale distributed archives presents severe bottlenecks [8, 13]. First, acquisition heterogeneity—arising from distinct detector geometries (secondary electron vs. backscattered electron), accelerating voltages, and working distances—induces substantial feature shifts in visual embeddings, causing visually dissimilar micrographs of identical metallurgical structures [7, 12]. Second, raw acquisition archives are frequently corrupted by operational artifacts including beam charging, illumination gradients, and sample contamination, which can silently degrade automated image analysis [18, 19]. Third, standard deep learning classifiers produce overconfident predictions on corrupted or out-of-distribution micrographs without providing provenance or comparative visual evidence [21].

To resolve these interconnected challenges, we present SCI-INTEL, an open, reproducible platform for acquisition-aware scientific image retrieval and quality-aware curation. Our core contributions are: (1) demonstrating a 66.23% reduction in the acquisition-geometry gap via representation adaptation; (2) uncovering an empirical trade-off between acquisition alignment and fine-grained artifact sensitivity; (3) introducing a modular dual-representation architecture combining DINOv2 and adapted projections without learned fusion; (4) implementing uncertainty-aware abstention and spatial localization; and (5) providing an operational evidence retrieval layer returning comparative micrographs for 100% of an evaluated $N=55$ query cohort.

Figure 1: Modular Dual-Representation SCI-INTEL System Architecture




Fig. 1. Modular Dual-Representation SCI-INTEL System Architecture. DINOv2 is dedicated to quality screening and localization; Phase-4 adaptation executes acquisition-aware retrieval.

II. RELATED WORK

A. Visual Representation Learning: Self-supervised Vision Transformers (DINO, DINOv2) [1, 2] learn rich representations across natural imagery, but lack intrinsic modeling of physical electron optical parameters.

B. Microstructure Retrieval: Prior works relied on static descriptors or ImageNet pre-training [12, 13], which degrade sharply under instrument transfer.

C. Image Quality & Uncertainty: No-reference image quality metrics [18, 19] combined with selective prediction [21] provide conservative review routing for ambiguous or corrupted scientific data.

III. SYSTEM ARCHITECTURE & METHODOLOGY

The SCI-INTEL architecture comprises: (1) metadata-aware ingestion assigning immutable SHA-256 digests; (2) dual representation extraction using frozen DINOv2 ViT-S/14 (for quality screening) and Phase-4 adapted projections (for retrieval); (3) quality-risk screening across 11 artifact classes; (4) spatial suspicious-region localization using patch feature residuals; (5) uncertainty-aware abstention using normalized Shannon entropy; and (6) evidence aggregation with deterministic lexical ranking.

IV. EXPERIMENTAL PROTOCOL & REPRODUCIBILITY

Evaluation spans 6,085 active micrographs: HCCI SEM (427 train, 135 val, 212 test), Carinthia SEM (4,591), and BBBC021 optical (720). The train, validation, and test partitions contain identical declared alloy specimen classes; hence, unseen-specimen generalization is not established. We enforce strict separation between Protocol M (historical masked exclusion) and Protocol U (authoritative unmasked distractors).

V. EXPERIMENTAL RESULTS

Representation	R@1	R@5	MRR	AUROC	Macro F1	Assigned Role
Frozen DINOv2 ViT-S/14	0.1321	0.9858	0.5200	0.8582	0.6837	Quality & Artifact Screening
Phase-4 Seed 42	0.1321	0.9953	0.5230	0.8251	0.6335	Acquisition-Aware Retrieval
Phase-4 Seed 123	0.1462	0.9906	0.5214	0.8215	0.6310	Acquisition-Aware Retrieval
Phase-4 Seed 2024	0.1557	0.9906	0.5338	0.8224	0.6324	Acquisition-Aware Retrieval
Phase-4 3-Seed Mean	0.1447	0.9921	0.5261	0.8230	0.6323	Acquisition-Aware Retrieval

TABLE I: Representation Comparison & Trade-off under Protocol U (Unmasked Distractors).




Fig. 2. Acquisition-Geometry Similarity Gap Reduction under Protocol U (66.23% mean gap reduction, $p = 5.03\text{e-}36$).

- A. Acquisition Robustness:* Phase-4 adaptation reduces the observed acquisition gap from 0.2016 to 0.0681 (66.23% reduction, $p = 5.03\text{e-}36$, $d_z = 2.19$), achieving Recall@5 of 0.9921 under Protocol U.
- B. Quality Screening Trade-off:* In contrast, frozen DINOv2 outperforms adapted projections by +5.14% Macro F1 on controlled artifact screening (0.6837 vs. 0.6323), demonstrating complementary specialization.
- C. Spatial Localization:* Saliency envelopes achieve mean IoU of 0.4454 across 500 test images.
- D. Operational Evidence Availability:* Across the evaluated $N=55$ cohort, comparative evidence was retrieved for 100.0% of queries with zero duplicates.
- E. Latency:* Pipeline serial execution averaged 23.40 ms/image (P95: 28.30 ms).

VI. DISCUSSION & DUAL COMPOSITION

Adapting representations to acquisition invariance projects out high-frequency sensor noise, which attenuates sensitivity to subtle image perturbations. A modular dual-representation architecture resolves this tension by assigning representations to their strongest evaluated roles without learning unverified fusion weights.

VII. SCIENTIFIC LIMITATIONS & BOUNDARIES

1. Train, validation, and test partitions contain identical declared alloy classes; unseen-specimen generalization is not established.
2. Controlled synthetic perturbations do not encompass all real-world physical microscope defects.
3. Human expert validation of scientific interpretation was not performed.
4. Evidence availability is an operational metric and does not establish visual correctness.
5. Evidence retrieval establishes geometric similarity; it does not constitute causal or diagnostic explanation.
6. Saliency bounding envelopes are model-derived suspicious regions, not confirmed physical defect boundaries.
7. OOD screening does not constitute open-world anomaly discovery.
8. Acquisition robustness is bounded to evaluated HCCI SEM conditions.
9. Absent metadata fields can limit evidence interpretation.
10. Measured latencies depend on the benchmark execution environment.
11. Protocol-M and Protocol-U retrieval metrics are not directly interchangeable.
12. The dual representation is a deterministic composition, not a learned fusion model.
13. No clinical diagnosis or medical decision-making claim is made.

VIII. CONCLUSION

SCI-INTEL provides a reproducible, scientifically verified platform for acquisition-aware scientific image retrieval and quality-aware curation.

REFERENCES

- [1] M. Oquab et al., 'DINOv2: Learning Robust Visual Features without Supervision,' TMLR, 2023.
- [2] M. Caron et al., 'Emerging Properties in Self-Supervised Vision Transformers,' in ICCV, 2021.
- [3] A. Dosovitskiy et al., 'An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale,' in ICLR, 2021.
- [4] K. He et al., 'Masked Autoencoders Are Scalable Vision Learners,' in CVPR, 2022.
- [5] P. Khosla et al., 'Supervised Contrastive Learning,' in NeurIPS, 2020.
- [6] T. Chen et al., 'A Simple Framework for Contrastive Learning of Visual Representations,' in ICML, 2020.
- [7] Y. Ganin et al., 'Domain-Adversarial Training of Neural Networks,' JMLR, 2016.
- [8] J. Johnson, M. Douze, and H. Jégou, 'Billion-Scale Similarity Search with GPUs,' IEEE Trans. Big Data, 2019.
- [9] M. Douze et al., 'The Faiss Library,' IEEE TPAMI, 2024.
- [10] Y. A. Malkov and D. A. Yashunin, 'Efficient and Robust Approximate Nearest Neighbor Search Using HNSW Graphs,' IEEE TPAMI, 2018.
- [11] H. Jégou, M. Douze, and C. Schmid, 'Product Quantization for Nearest Neighbor Search,' IEEE TPAMI, 2011.
- [12] B. L. DeCost, T. Francis, and E. A. Holm, 'Exploring the Microstructure Manifold: Image Representation, Similarity, and Retrieval in Materials Science,' IMMI, 2017.
- [13] J. Stuckner, B. Harder, and T. M. Smith, 'Microstructure Classification and Retrieval Using Computer Vision,' Comput. Mater. Sci., 2022.
- [14] K. Choudhary et al., 'Recent Advances and Applications of Deep Learning in Materials Science,' npj Comput. Mater., 2022.
- [15] T. Baltrušaitis, C. Ahuja, and L. P. Morency, 'Multimodal Machine Learning: A Survey and Taxonomy,' IEEE TPAMI, 2018.
- [16] A. Radford et al., 'Learning Transferable Visual Models From Natural Language Supervision,' in ICML, 2021.
- [17] R. J. Chen et al., 'Multimodal Co-Attention Transformer for Survival Prediction in Gigapixel Whole Slide Images,' IEEE TMI, 2021.
- [18] Z. Wang et al., 'Image Quality Assessment: From Error Visibility to Structural Similarity,' IEEE TIP, 2004.
- [19] A. Mittal, A. K. Moorthy, and A. C. Bovik, 'No-Reference Image Quality Assessment in the Spatial Domain,' IEEE TIP, 2012.
- [20] E. Krotkov, 'Focusing,' IJCV, 1987.
- [21] J. Yang, K. Zhou, Y. Li, and Z. Liu, 'Generalized Out-of-Distribution Detection: A Survey,' arXiv:2110.11334, 2021.
- [22] M. D. Wilkinson et al., 'The FAIR Guiding Principles for Scientific Data Management and Stewardship,' Scientific Data, 2016.
- [23] A. Paszke et al., 'PyTorch: An Imperative Style, High-Performance Deep Learning Library,' in NeurIPS, 2019.
- [24] C. Zauner, 'Implementation and Benchmarking of Perceptual Image Hash Functions,' Master thesis, FH Hagenberg, 2010.